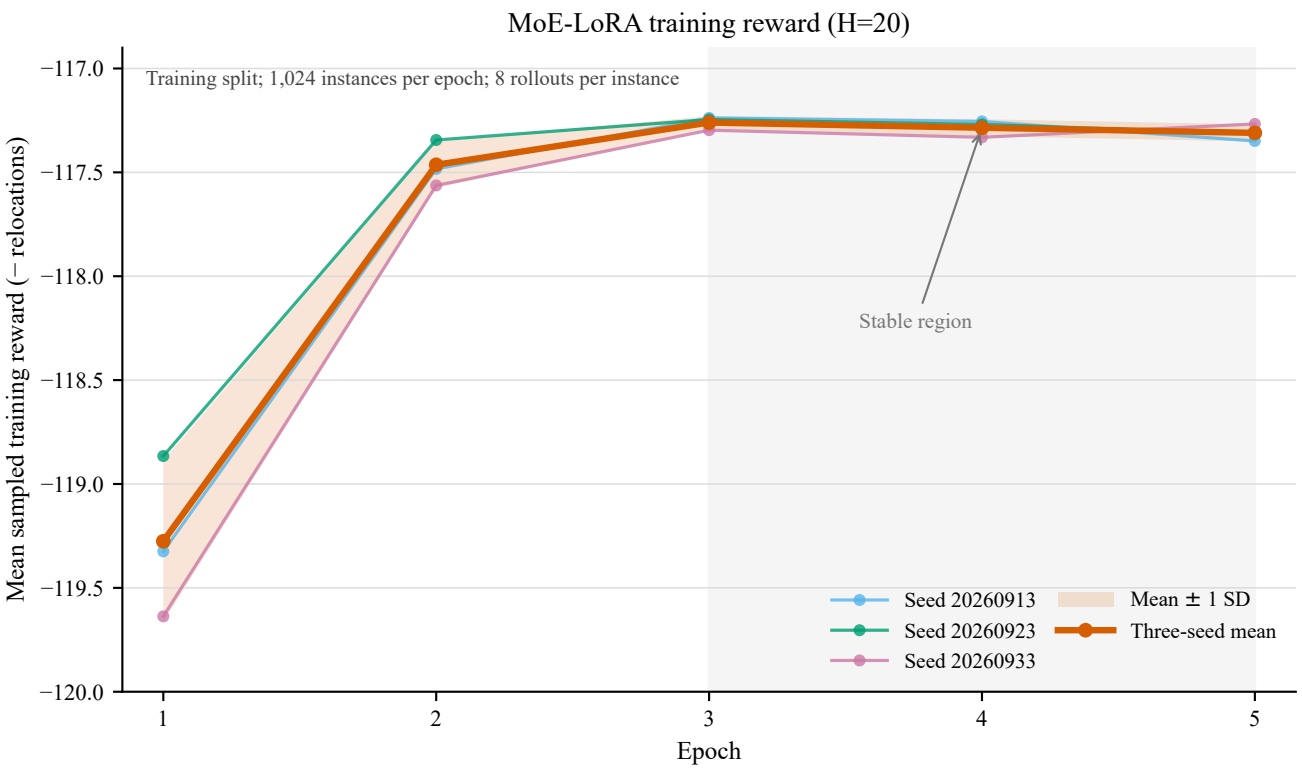
MoE-LoRA training reward (H=20)
Training split; 1,024 instances per epoch; 8 rollouts per instance
Mean sampled training reward (− relocations)
−117.0
−117.5
−118.0
−118.5
−119.0
−119.5
−120.0
Stable region
Seed 20260913
Seed 20260923
Seed 20260933
Mean ± 1 SD
Three-seed mean
1
2
3
4
5
Epoch